


André



DuckDB



DATA
SATURDAY
Holland



'98



(@ko_ver)



@andrekamman | [linkedin.com/in/andrekamman](https://www.linkedin.com/in/andrekamman) | andrekamman.com

Agenda

- What is it and where does it come from
- How do you install it and start it
- What makes it cool and different
 - Storage internals
 - FriendlySQL
 - Dataframe SQL
 - Formats and Extensions
 - CSV auto detection & Faulty line detection and handling
 - duck-dbt

What is it



- C++
- In-Process
- MIT License
- 15ms startup time
- CLI, Python, R, Rust, Go, Node.js, Java, C, C++, etc.

Where does it come from?

- **Mark Raasveldt & Hannes Mühleisen**
- **Started in 2019 while working at CWI in Amsterdam**
- **Dutch National center for Math and Computer Science**
- **Guido van Rossum created Python at the same place many years ago**
- **Hannes had a pet duck, hence “DuckDB”**



Version	<div><div>0.10.2 (stable release)</div><div>Nightly build (bleeding edge)</div></div>
Environment	<div><div>Command line</div><div>Python</div><div>R</div><div>Java</div><div>Node.js</div><div>Rust</div><div>Go</div><div>C/C++</div><div>ODBC</div></div>
Installation	<div><pre>pip install duckdb --upgrade</pre><div></div><p>On Windows, the DuckDB Python package requires the Microsoft Visual C++ Redistributable.</p></div>
Usage example	<div><pre>import duckdb cursor = duckdb.connect() print(cursor.execute('SELECT 42').fetchall())</pre><div></div></div>

Version	<div><div>0.10.2 (stable release)</div><div>Nightly build (bleeding edge)</div></div>
Environment	<div><div>Command line</div><div>Python</div><div>R</div><div>Java</div><div>Node.js</div><div>Rust</div><div>Go</div><div>C/C++</div><div>ODBC</div></div>
Platform	<div><div>Windows</div><div>macOS</div><div>Linux</div></div>
Download method	<div><div>Package manager</div><div>Direct download</div></div>
Installation	<div><pre>brew install duckdb</pre></div>
Usage example	<div><pre>./duckdb</pre><p><i># Note: DuckDB clients are installed without relying on any other DuckDB clients.</i></p><p><i># For example, the Python library can use a different version than the CLI client.</i></p><p><i># Therefore, they need to be updated separately.</i></p></div> 

Source: <https://duckdb.org/docs/installation>





Uncompressed

W

W

W

W

W

W

W

W

W

Constant

W

Uncompressed

W

W

W

A

A

X

X

X

X

Run-Length Encoding

W

A

X

3

2

4

Uncompressed

00000027

00000032

00000004

Bit Packing

27

32

04

Uncompressed

10000027

10000032

10000000

Frame of Reference

10000000

27

32

00

Uncompressed

Duck

Duck

Goose

Duck

Duck

Dictionary

0

0

1

0

0

Duck

Goose

Uncompressed

www.google.com

www.github.com

FSST

0

2

1

0

3

1

Symbol Table

www.

.com

google

github

- Chimp, based on Gorilla compression for floating points
- Patas takes that even further
- Compression does not apply to in-memory databases!!

Compression info via pragma_storage_info('<tablename>') function

<i>row_group_id</i>	<i>column_name</i>	<i>column_id</i>	<i>segment_type</i>	<i>count</i>	<i>compression</i>
4	extra	13	FLOAT	65536	Chimp
20	tip_amount	15	FLOAT	65536	Chimp
26	pickup_latitude	6	VALIDITY	65536	Constant
46	tolls_amount	16	FLOAT	65536	RLE
73	store_and_fwd_flag	8	VALIDITY	65536	Uncompressed
96	total_amount	17	VALIDITY	65536	Constant
111	total_amount	17	VALIDITY	65536	Constant
141	pickup_at	1	TIMESTAMP	52224	BitPacking
201	pickup_longitude	5	VALIDITY	65536	Constant
209	passenger_count	3	TINYINT	65536	BitPacking



Microsoft Windows [Version 10.0.22631.3447]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Andre>duckdb

v0.10.2 1601d94f94

Enter ".help" for usage hints.

Connected to a **transient in-memory database**.

Use ".open FILENAME" to reopen on a persistent database.

D INSTALL httpfs;

D LOAD httpfs;

D

D CREATE TABLE trek_facts AS

TM SELECT *

TM FROM 'https://blobs.duckdb.org/data/Star_Trek-Season_1.csv';

D

D DESCRIBE trek_facts;

column_name varchar	column_type varchar	null varchar	key varchar	default varchar	extra varchar
season_num	BIGINT	YES			
episode_num	BIGINT	YES			
aired_date	DATE	YES			
cnt_kirk_hookups	BIGINT	YES			
cnt_downed_redshirts	BIGINT	YES			
bool.aliens.almost.took.over.planet	BIGINT	YES			
bool.aliens.almost.took.over.enterprise	BIGINT	YES			
cnt_vulcan_nerve_pinch	BIGINT	YES			
cnt_warp_speed_orders	BIGINT	YES			
highest_warp_speed_issued	BIGINT	YES			


```
SELECT
```

```
    x_wing,  
    proton_torpedoes,  
    --targeting_computer
```

```
FROM luke_whats_wrong
```

```
GROUP BY x_wing, proton_torpedoes,
```

```
;
```

SELECT

```
'Use the Force, Luke'[:13] AS sliced_quote_1,  
substr('I am your father', 1, 4) AS sliced_quote_2,  
substring('Obi-Wan Kenobi, you''re my only hope',17,100) AS  
sliced_quote_3;
```

```
SELECT
    *
FROM (
    SELECT
        s1.tie_fighter,
        s2.tie_fighter
    FROM squadron_one s1
    JOIN squadron_two s2 ON 1=1 ) theyre_coming_in_too_fast;
```

tie_fighter

tie_fighter:1

green_one

green_two

SELECT

```
'These are the voyages of the starship Enterprise...' AS intro,  
instr(intro, 'starship') AS starship_loc,  
substr(intro, starship_loc + len('starship') + 1) AS trimmed_intro;
```

```
SELECT * EXCLUDE (jar_jar_binks, midichlorians) FROM star_wars;
```

```
SELECT
```

```
    sw.* EXCLUDE (jar_jar_binks, midichlorians),  
    ff.* EXCLUDE cancellation
```

```
FROM star_wars sw, firefly ff;
```

```
SELECT * REPLACE (
```

```
    movie_count+3 AS movie_count,  
    show_count*1000 AS show_count  
)
```

```
FROM star_wars_owned_by_disney;
```

SELECT

systems,

planets,

cities,

cantinas,

SUM(scum + villainy) AS total_scum_and_villainy

FROM star_wars_locations

GROUP BY ALL;

SELECT

*** EXCLUDE (cantinas, booths, scum, villainy),**

SUM(scum + villainy) AS total_scum_and_villainy

FROM star_wars_locations

GROUP BY ALL;


```
SELECT
    age,
    sum(civility) AS total_civility
FROM star_wars_universe
GROUP BY ALL
ORDER BY ALL;
```

It also supports DESC

And you can say NULLS FIRST or NULLS LAST

```
SELECT
    episode_num,
    COLUMNS('.*warp.*')
FROM trek_facts;
```

episode_num	cnt_warp_speed_orders	highest_warp_speed_issued
0	1	1
1	0	0
2	1	1
3	1	0
...
27	1	1
28	0	0
29	2	8

```
SELECT  
    MAX(COLUMNS('.*warp.*'))  
FROM trek_facts;
```

max(trek_facts.cnt_warp_speed_orders)

5

max(trek_facts.highest_warp_speed_issued)

8

```
SELECT
    episode_num,
    COLUMNS('.*warp.*')
FROM trek_facts
WHERE COLUMNS('.*warp.*') >= 2;
```

episode_num	cnt_warp_speed_orders	highest_warp_speed_issued
14	3	7
17	2	7
18	2	8
29	2	8

```
SELECT
    MAX(COLUMNS(* EXCLUDE season_num))
FROM trek_facts;
```

max(trek_facts. episode_num)	max(trek_facts. aired_date)	max(trek_facts. cnt_kirk_hookups)	...	max(trek_facts. bool_enterprise_s aved_the_day)
29	1967-04-13	2	...	1

```
SELECT
    MAX(COLUMNS(* REPLACE aired_date::timestamp AS aired_date))
FROM trek_facts;
```

max(trek_facts. season_num)	max(trek_facts. episode_num)	max(aired_date := CAST(aired_date AS TIMESTAMP))	...	max(trek_facts. bool_enterprise_s aved_the_day)
1	29	1967-04-13 00:00:00	...	1


```
SELECT
    episode_num,
    COLUMNS(col -> col LIKE '%warp%')
FROM trek_facts
WHERE COLUMNS(col -> col LIKE '%warp%') >= 2;
```

episode_num	cnt_warp_speed_orders	highest_warp_speed_issued
14	3	7
17	2	7
18	2	8
29	2	8

```
FROM my_table SELECT my_column;
```

```
FROM my_table;
```

```
COPY (FROM trek_facts) TO  
'phaser_filled_facts.parquet';
```

(OLD)

```
SELECT
  concat(
    list_aggr(
      string_split(
        UPPER('Make it stop'),
        ' '),
    'string_agg', '.'),
  '.' ) AS oof;
```

oof

MAKE.IT.STOP.

SELECT

('Make it so')

.UPPER()

.string_split(' ')

.list_aggr('string_agg','.'))

.concat('.') AS im_not_messing_around_number_one;

im_not_messing_around_number_one

MAKE.IT.SO.

```
CREATE TABLE proverbs AS
SELECT
    'Revenge is a dish best served cold' AS klingon_proverb
UNION ALL BY NAME
SELECT
    'You will be assimilated' AS borg_proverb,
    'If winning is not important, why keep score?' AS klingon_proverb;
FROM proverbs;
```

klingon_proverb	borg_proverb
Revenge is a dish best served cold	NULL
If winning is not important, why keep score?	You will be assimilated

```
INSERT INTO proverbs BY NAME
SELECT 'Resistance is futile' AS borg_proverb;
```

klingson_proverb	borg_proverb
Revenge is a dish best served cold	NULL
If winning is not important, why keep score?	You will be assimilated
NULL	Resistance is futile

item	year	count
phasers	2155	1035
phasers	2156	25039
phasers	2157	95000
photon torpedoes	2155	255
photon torpedoes	2156	17899
photon torpedoes	2157	87492

PIVOT purchases ON year
 USING SUM(count)
 GROUP BY item;

item	2155	2156	2157
phasers	1035	25039	95000
photon torpedoes	255	17899	87492

item	2155	2156	2157
phasers	1035	25039	95000
photon torpedoes	255	17899	87492

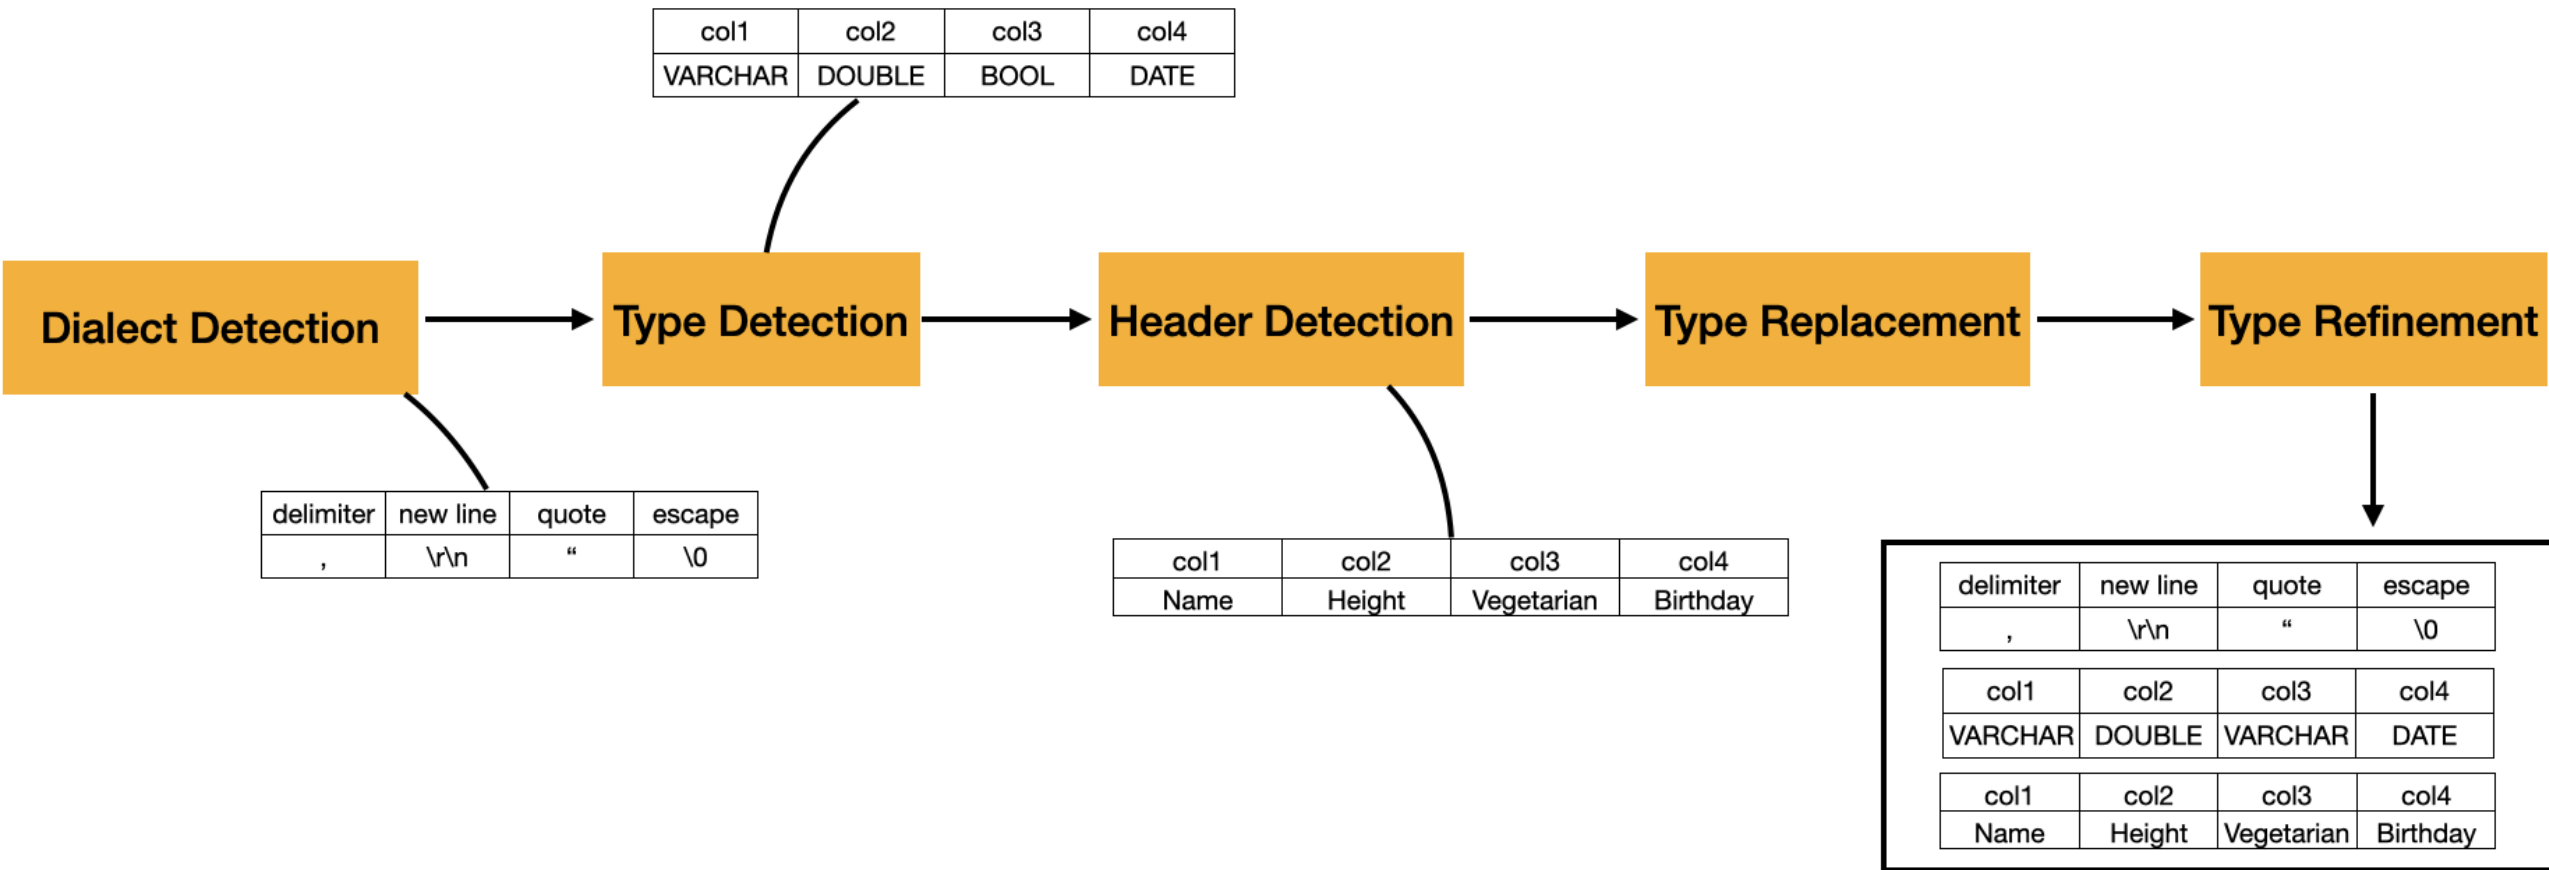
```
UNPIVOT pivoted_purchases
ON COLUMNS(* EXCLUDE item)
INTO
    NAME year
    VALUE count;
```

item	year	count
phasers	2155	1035
phasers	2156	25039
phasers	2157	95000
photon torpedoes	2155	255
photon torpedoes	2156	17899
photon torpedoes	2157	87492

And many many more!

- Number generator
- As-of joins
- Tpch data generator
- Json structs





Name, Height, Vegetarian, Birthday
"Pedro", 1.73, False, 30-07-92
... imagine 2048 consistent rows ...
"Mark", 1.72, N/A, 20-09-92

Conversion Error: CSV Error on Line: 5648

Original Line: Pedro,The 90s

Error when converting column "birth_date". date field value out of range: "The 90s", expected format

Column date is being converted as type DATE

This type was auto-detected from the CSV file.

Possible solutions:

- * Override the type for this column manually by setting the type explicitly, e.g. types={'birth_date': 'DATE'}
- * Set the sample size to a larger value to enable the auto-detection to scan more values, e.g. sample_size=20480
- * Use a COPY statement to automatically derive types from an existing table.

```
file= people.csv
delimiter = , (Auto-Detected)
quote = " (Auto-Detected)
escape = " (Auto-Detected)
new_line = \r\n (Auto-Detected)
header = true (Auto-Detected)
skip_rows = 0 (Auto-Detected)
date_format = (DD-MM-YYYY) (Auto-Detected)
timestamp_format = (Auto-Detected)
null_padding=0
sample_size=20480
ignore_errors=false
all_varchar=0
```

Name	Description	Type	Default
store_rejects	If set to true, any errors in the file will be skipped and stored in the default rejects temporary tables.	BOOLEAN	False
rejects_scan	Name of a temporary table where the information of the scan information of faulty CSV file are stored.	VARCHAR	reject_scans
rejects_table	Name of a temporary table where the information of the faulty lines of a CSV file are stored.	VARCHAR	reject_errors
rejects_limit	Upper limit on the number of faulty records from a CSV file that will be recorded in the rejects table. 0 is used when no limit should be applied.	BIGINT	0

scan_id	file_id	file_path	delimiter	quote	escape	newline_delimiter
5	0	faulty.csv	,	"	"	\n

line	line_byte_position	byte_position	column_idx	column_name	error_type	csv_line	error_message
2	10	23	2	age	CAST	Oogie Boogie, three	Error when converting column "age". Could not convert string " three" to 'INTEGER'



pythonTM

```
import duckdb
duckdb.sql("SELECT 42").show()
```

```
import duckdb
duckdb.read_csv("example.csv")
duckdb.read_parquet("example.parquet")
duckdb.read_json("example.json")

duckdb.sql("SELECT * FROM 'example.csv'")
duckdb.sql("SELECT * FROM 'example.parquet'")
duckdb.sql("SELECT * FROM 'example.json'")
```



```
import duckdb
```

```
# directly query a Pandas DataFrame
```

```
import pandas as pd
```

```
pandas_df = pd.DataFrame({"a": [42]})
```

```
duckdb.sql("SELECT * FROM pandas_df")
```

```
# directly query a Polars DataFrame
```

```
import polars as pl
```

```
polars_df = pl.DataFrame({"a": [42]})
```

```
duckdb.sql("SELECT * FROM polars_df")
```

```
# directly query a pyarrow table
```

```
import pyarrow as pa
```

```
arrow_table = pa.Table.from_pydict({"a": [42]})
```

```
duckdb.sql("SELECT * FROM arrow_table")
```

```
import duckdb
duckdb.sql("SELECT 42").fetchall()      # Python objects
duckdb.sql("SELECT 42").df()             # Pandas DataFrame
duckdb.sql("SELECT 42").pl()             # Polars DataFrame
duckdb.sql("SELECT 42").arrow()          # Arrow Table
duckdb.sql("SELECT 42").fetchnumpy()    # NumPy Arrays
```

```
import duckdb
duckdb.sql("SELECT 42").write_parquet("out.parquet") # Write to a Parquet file
duckdb.sql("SELECT 42").write_csv("out.csv")         # Write to a CSV file
duckdb.sql("COPY (SELECT 42) TO 'out.parquet'")      # Copy to a Parquet file
```



dbt